

# week5\_ipynb

February 2, 2021

## 1 Week 3: Classification

In this section we will apply what we have learned about the logistic regression model to fit a model and make predictions. We will be using the `penguins` dataset from `seaborn` and try to predict where or not a penguin is of Adelie species.

We will build a Logistic Regression model from scratch.

```
[55]: # Code for Week 5
import pandas as pd
import numpy as np
import seaborn as sns

# Import penguins
penguins = (sns.load_dataset("penguins")).dropna()
penguins["One"] = 1
penguins["Adelie"] = 1*(penguins["species"] == "Adelie")

print("Data Shape, ",penguins.shape)

# Take a look at the columns
print(penguins.head())

# What percentage of our data is Adelie
print(np.mean(penguins.Adelie))

# What features do we have
print(penguins.columns)
```

Data Shape, (333, 9)

	species	island	bill_length_mm	bill_depth_mm	flipper_length_mm	\
0	Adelie	Torgersen	39.1	18.7	181.0	
1	Adelie	Torgersen	39.5	17.4	186.0	
2	Adelie	Torgersen	40.3	18.0	195.0	
4	Adelie	Torgersen	36.7	19.3	193.0	
5	Adelie	Torgersen	39.3	20.6	190.0	

	body_mass_g	sex	One	Adelie
0	3750.0	MALE	1	1

```

1      3800.0  FEMALE    1      1
2      3250.0  FEMALE    1      1
4      3450.0  FEMALE    1      1
5      3650.0   MALE     1      1
0.43843843843843844
Index(['species', 'island', 'bill_length_mm', 'bill_depth_mm',
      'flipper_length_mm', 'body_mass_g', 'sex', 'One', 'Adelie'],
      dtype='object')

```

## 1.1 Setting up our Model

In this example, we will use the features “bill\_length\_mm”, “bill\_depth\_mm”, “flipper\_length\_mm”, and “body\_mass\_g” to predict whether or not the species is Adelie.

```

[56]: X = penguins[['One', 'bill_length_mm',
→ 'bill_depth_mm', 'flipper_length_mm', 'body_mass_g']]
Y = penguins['Adelie']

```

Recall that in logistic regression, we model the probability as

$$\pi(\mathbf{X}_i; \boldsymbol{\beta}) = \frac{1}{1 + \exp(-\beta_0 - \beta_1 X_{i,1} - \dots - \beta_p X_{i,p})}$$

To make predictions with this model (and evaluate the gradient) we will first need to write a function that takes in our feature matrix and a guessed value of  $\boldsymbol{\beta}$  and returns a vector of probabilities.

```

[57]: beta_initial = np.array((-0.001, 0.001, 0.001, -0.001, 0.001))
def LogitReg(Xtrain, beta):
    power = -1*np.dot(Xtrain, beta)
    pHat = 1/(1 + np.exp(power))
    return pHat
test = LogitReg(X, beta_initial)
# print(test)

```

We will also want a function that takes in the predicted probabilities and returns the evaluates the log-likelihood:

$$\ell(\boldsymbol{\beta}) = \sum_{i=1}^n Y_i \ln \pi(\mathbf{X}_i; \boldsymbol{\beta}) + (1 - Y_i) \ln(1 - \pi(\mathbf{X}_i; \boldsymbol{\beta}))$$

```

[58]: def LogitLikelihood(Ytrain, pHat):
    return 1/(len(Ytrain))*np.sum(Ytrain*np.log(pHat) + (1 - Ytrain)*np.log(1-
→ pHat))

test2 = LogitLikelihood(Y, test)
print(test2)

```

```
-2.5221444759399567
```

We will find the parameters  $\beta_0, \beta_1, \dots, \beta_p$  to maximize the (simplified) log-likelihood:

$$\ell(\beta) = \sum_{i=1}^n \left[ \ln \left\{ 1 + e^{\beta \mathbf{X}_i} \right\} - Y_i \beta \mathbf{X}_i \right]$$

The gradient of  $\ell(\beta)$  is given:

$$\nabla \ell(\beta) = [\tilde{\pi}(\mathbf{X}; \beta) - \mathbf{Y}] \cdot \mathbf{X}$$

where:

$$\tilde{\pi}(\mathbf{X}; \beta) = (\pi(\mathbf{X}_1; \beta), \dots, \pi(\mathbf{X}_n; \beta))'$$

denotes our vector of predicted probabilities at guess  $\beta$

In order to implement this, we will need to write a gradient descent function. We can use what we have above:

```
[86]: def LogitGradientDescent(beta_initial, num_iterations, gamma, Xtrain, Ytrain):
    # Set up for gradient descent
    beta = beta_initial
    likelihoods = []
    # Do the gradient descent (updating each time)
    for i in range(num_iterations):
        # Find the vector of probabilities
        pHat = LogitReg(Xtrain, beta)
        # Evaluate the log-likelihood function
        likelihood = LogitLikelihood(Ytrain, pHat)
        # Add the likelihood to the list
        likelihoods.append(likelihood)
        # Compute the gradient
        grad = (1/len(Ytrain))*np.dot(Ytrain - pHat, X)
        # Update Beta
        beta = beta + gamma*grad
    # Compute the likelihood for the final value of beta
    pHat = LogitReg(Xtrain, beta)
    likelihood = LogitLikelihood(Ytrain, pHat)
    likelihoods.append(likelihood)
    # Return the last value of beta and the likelihoods
    return beta, np.array(likelihoods)

from sklearn.model_selection import train_test_split
Xtrain, Xtest, Ytrain, Ytest = train_test_split(X, Y, random_state = 0)

num_iterations = 50000
gamma = 0.0000001

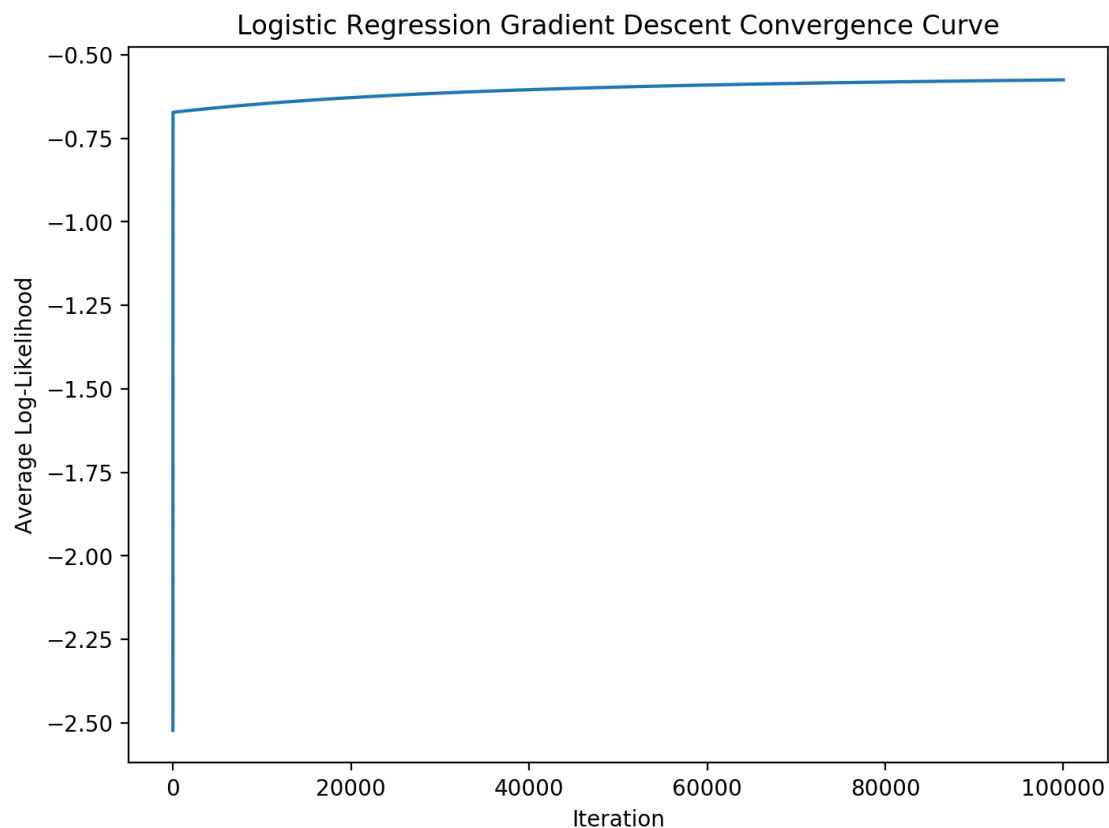
betaFinal, likelihoods = LogitGradientDescent(beta_initial, num_iterations,
↪gamma, Xtrain, Ytrain)
```

```
print(betaFinal)
```

```
[-0.00069094 -0.00481325  0.01057491  0.02554648 -0.00130612]
```

Now we plot the convergence curve:

```
[91]: import matplotlib.pyplot as plt
x = np.arange(len(likelihoods))
a = plt.figure(num=None, figsize=(8, 6), dpi=200, facecolor='w', edgecolor='k')
a = plt.title("Logistic Regression Gradient Descent Convergence Curve")
a = plt.ylabel("Average Log-Likelihood")
a = plt.xlabel("Iteration")
a = plt.plot(x, likelihoods)
a = plt.savefig("Convergence Curve")
```



And assess the performance of our model

```
[88]: pHatFinal = LogitReg(Xtest, betaFinal)
Yhat = 1*(pHatFinal >= 0.5)
accuracy = np.mean(Yhat == Ytest)
print(accuracy)
```

0.6428571428571429

How does this compare to the linear model?

```
[89]: from sklearn.linear_model import LinearRegression
      model = LinearRegression()
      model.fit(Xtrain,Ytrain)
      test = model.predict(Xtest)
      YHatLinear = 1*(test >= 0.5)
      linearAccuracy = np.mean(YHatLinear == Ytest)
      print(linearAccuracy)
```

0.9761904761904762

**Machine Learning is hard**